

From Lean Proofs to Public Claims

Release checks and trust boundaries in one formal mathematics repository

Will Cook

July 2026

Abstract

Lean verifies that a proof establishes the formal statement written in the source. It does not verify that a README or a paper describes that statement faithfully. Between a machine-checked theorem and a public mathematical claim there is a step no proof assistant performs: a person judges what the theorem means and what may responsibly be said about it. This paper describes how one public mathematics repository organises that step. A maintainer-reviewed claim record states, for selected results, the public wording, the supporting Lean declarations, the bounded domain, and the conclusion that remains open. A release checker then verifies recorded relationships: declaration names near recorded source lines, exact identity of the checked Lean source, required limitation wording, and freshly rebuilt generated files. Continuous integration runs that checker and the Lean build as two separate jobs. We follow one finite result through every layer and describe a deliberate exercise in which a false strengthening of the public wording passed every check because the relevant relationship had not yet been recorded; under a proved equivalence, the altered wording amounted to claiming an open problem settled. Recording the missing relationship afterwards makes the same false edit fail, which is coverage growth rather than demonstrated reliability. The checks preserve reviewed relationships; they do not interpret unrestricted prose, and the mathematical judgement remains human.

1 The publication gap

The repository at <https://github.com/wcook04/plectis-lean-erdos249-257> is a public Lean development around two unsolved problems in number theory, Erdős Problems 249 and 257. Both problems remain open. The project proves intermediate results, exact reformulations, and finite certificates around them; it does not claim a solution to either problem.

One theorem illustrates the problem this paper is about. Lean has checked a finite certificate at each of 28 listed scales, the largest labelled $t = 64$: at each listed scale, a finite calculation, checked in full by Lean, establishes the required property there. The registered open requirement asks for such certificates beyond every fixed cutoff. Whatever the largest checked scale is, a cutoff lies beyond it, and the finite list says nothing there. The finite list therefore does not settle the open problem.

A README can nevertheless drift, and the stakes are easy to state. By an equivalence proved in the development (Section 3), certificates beyond every fixed cutoff are not merely better evidence: their existence is equivalent to the irrationality of Problem 249 itself. Three objects therefore recur throughout this paper: the *finite theorem*, certificates at the 28 listed scales; the *open requirement*, certificates beyond every cutoff; and the *equivalence* between the open requirement and the irrationality. Lean has checked the first and the third; no Lean theorem carries the first to the second, and producing one would settle the problem. An edit that presents the finite cases as completing the open requirement therefore asserts, in English, exactly the bridge the development does not contain; in substance it announces a solution to an open problem. Every Lean proof remains valid under that edit, because the README is

not part of any proof. No program in the repository decides whether a line of English has the same meaning as a formal statement; a machine-checked theorem is, in that sense, not a self-publishing claim.

The repository keeps a *maintainer-reviewed claim record* in `docs/claims.json`. For each selected result the record states the public wording, its status, the Lean declarations that support it, the bounded domain it covers, and the registered conclusion that remains open. A release checker, `scripts/check_release.py`, then verifies that the recorded relationships still hold across the Lean source, the record itself, the authored public pages, and the generated files. The record covers *registered* claims, not every line of prose in the repository; software cannot preserve a relationship it was never pointed at, and Section 5 shows this limit operating in practice. The shortest accurate summary of the division of labour is:

Lean checks the formal proofs. A maintainer reviews what those proofs mean and how they may be described. The release machinery checks that the recorded relationships remain intact before a release.

2 The release workflow




Figure 1: The configured checks for one revision. Dashed elements are human judgements; solid elements are files and programs. Three relations of different kinds meet here. Lean establishes the *proof relation*: the source proves the stated theorems. Maintainer review asserts the *semantic relation*: the recorded wording describes those theorems within stated limits. The release job checks *persistence relations*: recorded names, wording, identities, and generated files still agree. Continuous integration reruns the first and third; nothing performs the second mechanically, and neither job reads unrestricted prose. In the worked example the protected object is the recorded boundary between the finite theorem and the open requirement.

Figure 1 shows the path in five visible parts: the Lean source; the maintainer-reviewed claim record; authored public documents; generated indexes and summaries derived from the record; and the release program and continuous-integration workflow that rerun the checks.

Establishing that public wording is faithful to a formal theorem, within stated limits, is a semantic act: someone must read both and judge. Preserving recorded consequences of that judgement is mechanical: a program can compare names, files, and wording against a checklist. The machinery performs only the second act, answering a narrow question: once a person has approved a proof-to-claim relationship, how are later revisions kept from silently eroding it? Whether the approval was correct is a different question, and no part of the design answers it.

A boundary of a different kind should be fixed equally early. The repository does not technically force a second independent mathematician to approve a change. A single maintainer can edit a Lean statement, its claim record, and the public wording together, and every check can pass while the shared interpretation is wrong; Section 6 instantiates this on the worked example. The checks are drift detection for reviewed decisions, not a substitute for review, and they add no authority of their own.

The two checks answer different questions. `lake build` asks whether Lean accepts the imported formal statements and proofs; it knows nothing about the README, and it would accept, unchanged, a revision whose README announces the open requirement met, because no public document is among its inputs. `python3 scripts/check_release.py` asks whether the recorded relationships among proofs, record, public pages, and generated files still hold; it confirms the identity of the Lean source but does not run Lean. The workflow in `.github/workflows/lean.yml` runs them as two separate jobs on every push and pull request, in a clean checkout. A failed comparison makes `scripts/check_release.py` exit with an error, and the configured workflow treats that exit as a failing job. Blocking is a property of the repository’s configuration, not of the script: the script can only refuse to succeed. The workflow’s power is operational, not mathematical: a failing job can stop a configured route to release; a passing one confers no mathematical standing.

The generated views deserve demystification. They select and reorganise recorded information (a reading map, a code index, a machine-readable summary) so that a reader or a program can enter the repository without opening every file. They create no new mathematics, each names the program that builds it, and a stale view is rebuilt by that owner rather than repaired by hand. The complete file map, ownership table, and maintenance commands live in [ARCHITECTURE.md](#); this paper keeps what the trust argument needs.

3 One claim from Lean theorem to public page

The architecture becomes concrete when one result is followed through every layer. We begin with its public meaning, not its internal name:

Lean has checked a finite certificate at each of 28 listed scales through $t = 64$. The registered open requirement asks for certificates beyond every fixed cutoff. The finite list does not settle it.

3.1 What the certificates are for

The result belongs to Erdős Problem 249, which asks whether the totient constant

$$S = \sum_{n \geq 1} \frac{\varphi(n)}{2^n}$$

is irrational, where $\varphi(n)$ counts the integers in $\{1, \dots, n\}$ sharing no factor greater than 1 with n . The certificates take their force from the equivalence, which is worth seeing at human scale.

Write R_N for the tail of the series beyond position N , rescaled by 2^N , so that R_N is $2^N S$ minus an integer. If S were a fraction, its binary digits would eventually repeat with some period h , and the development proves that this repetition forces the differences $R_{N+h} - R_N$ to be integers from some position on; conversely, a single integral difference at any positive shift already forces S to be rational. The chain checked in Lean is exact: S is irrational precisely when every such difference, over every period h and every position N , fails to be an integer; and, at fixed parameters, some finite residue computation with a rigorous tail allowance succeeds exactly when the failure holds. A *certificate* at fixed parameters is that finite computation, and it is integer arithmetic throughout: a finite weighted sum of totient values is an integer, its residue modulo a power of two is computed exactly, and a rigorous bound on the discarded infinite tail widens the forbidden value zero to a small neighbourhood; the certificate checks that the residue lies outside it. The project calls a kernel-checked certificate a *certified kill*, and the name marks an ordinary exclusion argument: each certificate eliminates the family of candidate rational representations of S tied to its parameters: the fractions that would have made that particular difference an integer.

The scales of the finite theorem compress this two-parameter requirement to one. The diagonal $H_t = \text{lcm}(1, \dots, t)$ is divisible by every candidate period up to t , so, roughly speaking, one certificate along the diagonal, which fixes both the period and the position to H_t , speaks to all of those periods at once. The development proves the compression exact: certificates along the diagonal beyond every fixed cutoff are equivalent to the full requirement, and therefore, through the chain above, to the irrationality itself. Certificates on a finite list, however long, are not.

Writing $\text{Cert}(t)$ for “a certificate exists at scale t ”, the two statements then read:

checked: $\text{Cert}(t)$ for every t in a fixed list of 28 scales;
open: for every cutoff T , $\text{Cert}(t)$ for some $t > T$.

The boundary between them is a matter of logical form, not of scale. Twenty-eight scales, two hundred and eighty, or twenty-eight thousand stand in the same relation to the open requirement: each is a fixed finite list with a cutoff beyond it. No finite list, however long or expensive to verify, implies the second statement without a further theorem that produces new witnesses. Thus $t = 64$ is the endpoint of this verified computation, not a frontier that Lean has proved will continue. This is the boundary an edit can silently delete.

3.2 The formal evidence

The Lean module `Erdos249257/DiagonalPincerCertificatesT64.lean` defines the explicit list of scales: the 28 values of t up to 64 at which the diagonal H_t attains a new value, which are the prime powers up to 64 together with the initial scale 1. The sparseness is structural, not a sampling choice: between prime powers the diagonal does not move, so no new statement exists to check. Its assembling theorem is `certifiedKill_diagonal_all_imported_through_t64`: for every listed t , the theorem gives a finite checking depth d_t and a decidable arithmetic witness at (H_t, H_t, d_t) : period H_t , position H_t , depth d_t . Imported modules prove the cases; the terminal theorem combines them. The proofs use kernel-checked arithmetic, including the tactic `decide`, which evaluates a decidable statement inside Lean’s proof kernel. The equivalence is likewise named formal evidence, not a promise of this paper. The module `Erdos249257/LcmConeFlatness.lean` proves the chain of Section 3.1: the tail-difference and pointwise certificate equivalences, and past them the diagonal characterisation `irrational_totient_series_iff_lcm_diagonal_certificate_supply`. The root `Erdos249257.lean` imports both modules, so `lake build` checks them.

3.3 The reviewed record

Table 1 shows the substance of the corresponding entry in `docs/claims.json`, ordered as a reader meets it: what is claimed, where it stops, what remains open, then the supporting machinery. The entry also records an exact source line for each declaration; those coordinates are omitted from the table.

Field	Content
Public statement	“Kernel-checked certified kills for small periods and at 28 explicit lcm-diagonal scales through $t = 64$, plus finite certificate shards, each discharged by decide .”
Bounded domain	The listed small periods, the 28 scales above, and the parameters encoded by the finite certificate shards.
Remaining open	A link to the registered open requirement <i>unbounded certificate supply</i> : produce certified witnesses at unbounded parameters.
Status	<i>verified finite instance</i> : Lean checked the stated finite inputs.
Supporting declarations	Five named Lean declarations with recorded modules and source lines, ending in <code>certifiedKill_diagonal_all_imported_through_t64</code> .
Claim identifier	<code>certified_kill_instances</code>

Table 1: The reviewed claim entry for the worked example, exact source coordinates omitted. The pairing of the positive statement with its bounded domain and its remaining-open link is the point: public overstatement usually happens by deleting a boundary, not by inventing a theorem.

The public statement is exact but not self-explanatory. Two of its terms appear in Section 3.1, *certified kills* and the *lcm-diagonal scales* H_t ; the *small periods* are the periods 1 through 8, each killed by one explicitly checked certificate; and *certificate shards* are smaller reusable certificates at fixed parameters, each discharged by the kernel tactic **decide**. The remaining-open field is the open requirement of Section 3.1, and under the equivalence it is part of the claim’s identity rather than an optional caution: remove it, and the entry no longer separates the finite theorem from an answer to Problem 249. A record can be precise without being pedagogy: its job is to be the stable object that explanatory prose, the release checker, and the generated views refer to; the explanation a cold reader needs is the job of prose like this section.

The fields differ in kind, and the difference matters for what “reviewed” means. The public statement, bounded domain, remaining-open link, and status are semantic fields: they record the outcome of a human judgement, and only a person can meaningfully change them; the status vocabulary is fixed in the record, and the checker rejects a status outside it. The claim identifier is a lookup handle for software, which is why it appears last rather than in place of the wording, and the per-declaration source lines are bookkeeping that tooling may refresh when code moves; refreshing them re-establishes nothing semantic. What no field can do is make the English faithful to the Lean: the maintainer judged that the public statement above describes the five declarations correctly, and the record stores the outcome of that judgement, not a justification of it. The judgement itself is mundane: read the terminal theorem, check that the wording claims the listed scales and no more, confirm that the unbounded requirement stays named as open. It is also an event rather than a permanent property: the review covered particular declarations and particular wording, so a later change to either reopens it, and the machinery of the next sections can only carry its recorded consequences forward in time. What that decision must survive is not a rival mathematician but an ordinary rewrite months later.

3.4 What the checker verifies here

For this claim, `scripts/check_release.py` verifies a finite, listed set of relationships. Each declaration name must appear in the Lean source within three lines of its recorded coordinate; this is a drift check on names and positions, not a check that the declarations entail the public wording. The checked Lean source itself must be exactly the formal-source checkpoint named in the claim record: the checker compares the supported root and the whole library directory against that recorded version and rejects the release if they differ or if untracked Lean files are present, so the verified names cannot refer to an older source than the one being shipped. The required limitation wording must remain present on the registered public pages; for this claim that wording is the visible trace of the boundary of Section 3.1, the clause keeping the open requirement open. The generated views must equal a fresh rebuild from their sources. Presence, identity, and proximity are all the checker can test. Lean checks the theorem; the checker checks the recorded description around it; a person remains responsible for the claim that the two say the same thing.

4 What the checks establish

Four kinds of assurance meet in Figure 1, and informal talk about “verification” blends them. The figure’s relations say what is asserted; an assurance says why it may be accepted; execution, in particular, grounds a claim about an event, not a further relation between artefacts. Lean provides *proof* assurance: the written formal statements follow from their imported environment. Maintainer review provides *interpretive* assurance: a recorded judgement that selected wording describes selected theorems within stated limits. The release checker provides *persistence* assurance: the recorded relationships still hold at this revision. Continuous integration provides *execution* assurance: the configured jobs ran on the named revision and passed. These are not interchangeable, and success at one level does not flow upward. A green run establishes a conjunction: Lean accepted the proofs, and every recorded relationship held. It does not establish that the wording is faithful, nor that every consequential claim was registered. Table 2 states the guarantee and non-guarantee of each layer, because combining layers in prose makes the guarantee sound stronger than it is.

Instantiated on the worked example, the four assurances read as follows. Proof assurance: Lean accepts the finite theorem and the equivalence, and accepts nothing that carries the finite list to the open requirement. Interpretive assurance: a maintainer’s recorded judgement that the public wording asserts the finite theorem, confines it to the listed scales, and keeps the open requirement named as open. Persistence assurance: the recorded traces of that judgement (declaration names near their coordinates, exact source identity, the limitation clause on the public pages, freshly rebuilt views) still hold at the current revision. Execution assurance: both jobs ran on that revision in a clean checkout and passed. The non-guarantees follow the same lines. Lean cannot notice English that quietly implies the open requirement met. Review cannot promise that every page repeating the claim was found. The checker cannot reject a paraphrase it was never given. A green run cannot make a mistaken reading of the equivalence correct.

Parts of the release checker are stronger than the phrase “consistency check” suggests, and worth stating. As a project-source policy, the checker rejects `sorry`, `admit`, new `axiom` declarations, and `native_decide`; the last exclusion keeps finite certificates on the repository’s kernel-checkable computation path. This lexical policy does not show that imported libraries use no classical axioms, and the pinned toolchain and imported definitions remain part of the trust base. The checker also compares the entire public proof surface byte for byte against the recorded formal-source checkpoint and rejects untracked Lean files. It rebuilds the generated views and requires exact agreement rather than plausible freshness, verifies the recorded source and PDF fingerprints of the shipped papers, and carries its own negative examples: deliberately

Layer	What it can establish	What it cannot establish
Lean build and kernel	The written formal statements are proved from their imported definitions and assumptions, in the pinned toolchain.	That a formal statement is the one the author intended, or that any English description of it is faithful.
Maintainer review	A recorded judgement that a formal statement has the intended meaning and that selected public wording describes it within stated limits.	That Lean accepts the proofs; that every important claim was noticed and registered; its own independence or quality.
Release checker	That recorded names, coordinates, fields, links, required wording, source identity, and generated files agree with the declared rules; that the Lean source avoids the banned proof commands; and that its own deliberately wrong examples still fail.	The meaning of unregistered prose; the completeness of the recorded checklist; the correctness of the human judgements it preserves.
Continuous integration	That both jobs ran in a clean checkout of the pushed revision and exited successfully.	Independent mathematical approval; anything beyond what the two jobs themselves establish.

Table 2: The trust boundary, one layer at a time.

wrong inputs that must keep failing, so that a check cannot silently become vacuous. None of this reaches semantics. A fingerprint can prove two files identical; it cannot say which of two differing files expresses the right mathematics.

When a file changes, the revision passes through the same stations: proof validation with `lake build`; semantic reconsideration, a mathematical judgement about whether statements, assumptions, or intended meaning moved; record and page updates where public meaning changed, under the review rules in `docs/methodology.json`; regeneration of derived views by their named builders; and structural comparison with `python3 scripts/check_release.py`, after which the push repeats both jobs in a clean checkout. A prose-only edit skips only the first station: if the words now make a stronger or different mathematical claim, the record and the review must change with them. The command list is maintenance detail in the repository guide.

5 A boundary the checklist missed

The limits in Table 2 are not hypothetical. They were probed by a deliberate exercise, recorded in `docs/publication_evidence.json`, whose shape is worth describing exactly. The reporting register below is deliberate: the original run logs were not retained, so the account rests on the retained record, and Appendix A states what survives.

The retained maintainer record reports that, at a pinned repository checkpoint, identified exactly in the evidence file, ten deliberate false edits were applied to an isolated working copy, one at a time, with the copy restored between runs and the full release gate run after each edit (the release checker alone; it does not run Lean, so the formal proof job was not part of the per-edit loop). The edits targeted the seven recorded relationship families: status labels, record structure, source coordinates, paper link anchors, boundary wording, generated-view freshness, and size budgets. Representative edits upgraded a conditional status to “proved here”, deleted an open-boundary clause from the README, drifted a declaration coordinate by forty lines, introduced `native_decide` into a proof, and hand-edited a generated module count by one.

According to that record, nine of the ten edits were rejected. One escaped: an edit strength-

ening a finite-evidence boundary clause in the README, replacing wording that the finite cases *do not supply* the open requirement with wording that they *complete* it (both clauses paraphrased into this paper’s vocabulary; the reconstruction manifest of Appendix A retains the exact text). Under the registered equivalence of Section 3.1, the altered page claimed, in effect, that the open problem was settled: the edit wrote into the README exactly the bridge from the finite theorem to the open requirement that no Lean declaration supplies. The formal proofs were untouched, and the release gate passed in full, because that clause was not among the recorded relationships. The error was outside the written checklist, and the checker preserved exactly what it had been given.

One escape carries more information than nine rejections: it refutes, by counterexample, the strongest reading a green run invites, that passing every configured check certifies faithful public prose. The refutation needs no statistics, and the exercise supplies none.

The repair had two parts. The clause became a declared anchor: required wording on the public page, recorded in the checklist. And the checklist derives from each such anchor an opposing example: a copy of the page with the clause weakened or removed must be rejected. The intact text was then rechecked and passed; the same false edit now fails. The retained failing copy is a *boundary witness*: it demonstrates that, for this one relationship, the check separates the approved page from a known false neighbour, and it demonstrates nothing else; that is coverage growth, not general effectiveness. The other nine edits were not rerun against the extended checklist, and no aggregate result is claimed for the repaired configuration.

This exercise demonstrates a coverage boundary, not a reliability score. Its ceilings should be stated as plainly as its outcome. The edits were authored by the checker’s author, so the human review gate of Figure 1 was deliberately out of the loop and the exercise measures the recorded checklist alone; most relationship families were probed by exactly one edit. The exercise covers one repository, one maintainer, and one working style, with no comparison against ordinary continuous integration or careful manual review. The nine reported rejections rest on the retained record alone; only the intact baseline and the escaped edit have present-day replays. The original run logs were not retained; the evidence file records the protocol, the outcomes, and these limitations, and a deterministic reconstruction of the edits can be replayed against the current checker, but the reconstruction is not the original evidence. What survives all of these ceilings is the design lesson:

The program can preserve a relationship only after a person has identified and recorded that relationship.

The residual risk follows the same line: requiring a clause to be present does not detect a contradictory stronger claim elsewhere on the page. Coverage grows only by the route above: notice a relationship, record it, give it an example that must fail. The exercise is one full turn of that loop.

The shape of this limit is familiar from Section 3.1. The recorded checklist stands to unrestricted prose roughly as the finite certificate list stands to the open requirement: every registered relationship is real and checked, and no finite list of relationships exhausts an open-ended domain. The parallel is structural rather than exact, since prose enjoys no equivalence theorem and no diagonal compresses it, but it says why the checker’s silence outside its records is a boundary to be governed rather than a defect to be patched once.

6 Scope, reuse, and limits

The failures the design addresses. The checks are aimed at accidental, single-point drift after a sound review: a declaration moves and its recorded coordinate goes stale; a generated view is hand-edited or stale; a required limitation clause disappears in a rewrite; a status is

quietly upgraded; the shipped Lean tree ceases to be the reviewed one; a check decays into one that can no longer fail. Each changes one endpoint of a recorded relationship while the others stand still; mechanical comparison suits that shape of failure.

The failures it does not address. Three families remain outside by design. First, semantically equivalent damage: a paraphrase that strengthens a claim without touching a registered anchor, a contradictory assertion elsewhere on the page, misleading emphasis, or a new public document the record never mentions; requiring a clause to be present is not the same as requiring the page to agree with it. Second, coordinated change: a maintainer with authority over source, record, and prose can move all three together, and perfect structural agreement then preserves a shared error. Were the certificate notion itself weakened in the source while record and prose were updated to match, every comparison could agree while the revised reading of the equivalence is wrong; what remains is the quality and governance of the review itself. Third, wrong review: if the original judgement was mistaken, persistence faithfully preserves the mistake. The machinery has no opinion about the judgements it protects.

Essential and contingent limits. Some limits are essential to the chosen shape: while claims live in unrestricted prose beside an external record, human interpretation cannot be eliminated, and coverage is selective. The record draws an *assurance perimeter*: inside it, named commitments carry mechanical obligations; outside it, silence is expected rather than exceptional. Where the perimeter runs is itself a judgement, weighted here toward headline results and wording whose accidental strengthening would change the repository’s public status, as one deleted boundary clause would under the equivalence. Unregistered wording is invisible, and so is an unregistered copy: the record stores each claim once, a claim can appear on many pages, and the perimeter encloses listed appearances rather than meanings. Others are contingent and could be engineered away without changing the architecture: the single maintainer (one person here performs two separable acts, judging that wording is faithful and authorising it as the project’s wording), the three-line proximity window, exact-wording anchors, one edit per relationship family in the exercise, and the unretained logs.

Reuse. Another formal project could reuse the pattern without copying any file format, and the worked example says where the pattern pays: wherever bounded formal evidence sits beside a stronger public target. Finite instances beside a universal statement, as here; a conditional theorem beside its unresolved hypothesis, the shape of this repository’s conditional reductions; one direction of an implication beside a claimed equivalence; a result under named assumptions beside an unconditional headline. In each case the portable unit is the reviewed boundary between what is established and the adjacent stronger statement that is not. The portable obligations are short to state: a defined unit of public claim; each claim naming its formal evidence and the source state that evidence lives in; explicit scope, assumptions, or finite domain, with consequential non-conclusions recordable; a human review act distinguishable from machine validation; named owners for derived public views; proof checking and publication checking as separate checks, each able to fail the configured workflow; a deliberately disagreeing example for every invariant that matters; and coverage represented as selective rather than assumed complete. Stronger variants lie beyond this minimum, each buying coverage at a cost: generating the public wording from the record instead of anchoring authored prose, which would have denied the escaped edit of Section 5 its freedom; extracting statements and assumptions from the proof assistant instead of recording source lines; requiring independent review for designated claims; constraining critical claims to a controlled vocabulary; or adding advisory semantic tooling that proposes, but never certifies, candidate discrepancies. This repository occupies that family’s conservative end.

What is not established. The architecture does not establish a solution to either Erdős problem; that a formal statement is the statement the author intended; that software understood any unregistered prose; that the record is complete; that one maintainer’s review is independent or adequate; or that the design transfers to other projects with its behaviour intact. A large number of passing comparisons measures none of those things. What a green workflow run establishes is narrower: Lean accepted the formal proofs, and every configured comparison passed for the named revision. This paper is an architecture note with one bounded case study, not an empirical evaluation of the design. An evaluation would have to be built in stages, each with its own ceiling: authenticated records of future review and release runs, which establish provenance and denominators but no effectiveness; challenge edits authored by people other than the checker’s author, which measure response to unseen faults but not field behaviour; and only then comparison against ordinary review on the same changes. None of those stages exists yet, and this paper claims none of them.

7 Related systems

Related systems sit on two axes: where in the life of a formalisation they act, and where their semantic structure lives. Proof blueprints act before and during formalisation, connecting plans to Lean declarations and tracking dependency and completion: `leanblueprint` for authored blueprints [1], `LeanArchitect` for generating and maintaining them [2]. The claim record here acts after: not “what remains to formalise”, which for a blueprint would list the open requirement as planned work, but “what may the finished formalisation be said to establish, and with what limits”. Lean Atlas supports judging whether a formal statement expresses the intended mathematics [3], the gap between intention and formal statement; the record here takes that judgement as input and works one gap later, between formal statement and public claim. On the second axis, Isabelle/DOF places typed, checked structure inside formal documents themselves [4], gaining coverage over whatever is written in that structure at the cost of constraining the document; the bounded domain and remaining-open fields of Table 1 could live there as typed fields. This repository instead leaves prose unrestricted and checks an external record, buying a clear human review surface at the cost of open-world coverage. Documentation-consistency systems such as `CASCADE` interpret prose to generate tests [5], reaching text that explicit records cannot see at the price of probabilistic interpretation: such a system might have flagged the strengthened README of Section 5 unprompted, and might equally flag a faithful rewrite. The design here takes the opposite point and leaves unrecorded drift invisible. The deliberately wrong copies of Section 5 are known-answer regression fixtures in the mutation-analysis tradition [6, 7], keeping checks non-vacuous rather than computing a mutation score; the ten-edit exercise is itself a small hand-authored mutation run, and its stated ceilings, author-written edits, single probes, no rerun, are exactly the concerns that tradition studies. Outside formal mathematics, the claim record is a lightweight relative of requirements traceability [8] and of assurance cases [9]: a claim with named evidence, explicit scope, and recorded non-conclusions is close to a claim–evidence–warrant structure, except that the warrant here is an unformalised human judgement and the known defeaters, unregistered prose and reviewer error, are stated rather than argued away, a discipline assurance cases share through their explicitly undeveloped goals. What formal mathematics adds to that older picture is an unusually strong evidence layer beside an unusually mutable public one; the distance between them is the subject of this paper.

8 Conclusion

The paper opened with a distance: 28 checked scales on one side, certificates beyond every fixed cutoff on the other, and a proved equivalence making the far side the open problem itself.

Everything afterwards answers one question about that distance: what would it take for the words that mark it, the listed scales, the finite domain, the registered open requirement, to survive years of edits by the hands free to change them?

The answer is narrow and exact within its bounds. The workflow does not make prose formally true. It makes a selected set of proof-to-prose relationships explicit, reviewed, and mechanically persistent, and it makes the configured workflow fail when one of them breaks. The worked example shows both halves: the repaired checker accepts the intact wording and rejects the reconstructed false edit, while the earlier unrecorded boundary passed unnoticed. Three achievements should stay separate, because none implies the others: a theorem is proved; its public interpretation is reviewed; the reviewed relationship is preserved. Lean supplies the first. Only a person can supply the second. The machinery here supplies the third, for the relationships it has been given.

The whole case fits in one chain. Lean proves the finite theorem. Lean proves the equivalence. No theorem carries the one to the other. A person approved wording that keeps the distinction. An edit erased the distinction without touching a proof. Every configured check passed while the relationship was unrecorded. The relationship is now recorded and its known false form fails. Unrecorded false forms remain possible.

The evidence here should be weighed at three strengths. That the architecture exists and behaves as described at one revision is demonstrated. That reviewed proof-to-prose relationships can be made mechanically persistent, without pretending that software interpreted the prose, is argued, and the case study supports it. That such an architecture reduces consequential publication drift at acceptable maintenance cost across projects is a hypothesis for future evaluation, not a result of this paper. Both Erdős problems remain open, and nothing here changes their status: the finite theorem is proved, the equivalence is proved, and the open requirement, the unbounded certificate supply, is exactly as unproved after these checks as before them.

A Reproducibility

The reviewed claim record is `docs/claims.json`. Its release block names the saved Git revision of the Lean source, and the release checker requires an exact match. The paper omits that changing identifier because the record is its single owner; a copy in prose would be a second value that can drift.

The paper inventory is `docs/publication_contract.json`: source and rendered fingerprints for each shipped manuscript, with validation commands. The record of the Section 5 exercise is `docs/publication_evidence.json`: protocol, outcomes, and limitations. Its deterministic reconstruction is described by `experiments/publication_mutations.json`; the replay command is `scripts/run_publication_mutations.py` with `--verify-operators`. The original raw outputs of the historical runs were not retained, and the reconstruction does not claim to be them.

The papers build with Tectonic or standard $\text{\LaTeX}$ via `make -C paper`. Buildability, not byte-identical output across TeX engines, is the claim. The shipped PDF is the tracked root copy, whose fingerprint the paper inventory records and checks. These identities establish which artefacts are named, not that the artefacts are interpreted correctly; the appendix is one more instance of the paper’s subject.

References

- [1] P. Massot, *leanblueprint: a plasTeX plugin for formalization blueprints*, [project repository](#), accessed July 2026.
- [2] T. Zhu, P. Monticone, J. Avigad, and S. Welleck, *LeanArchitect: Automating Blueprint Generation for Humans and AI*, 2026, [arXiv:2601.22554](#).

- [3] B. Yanahama and A. Sannai, *Lean Atlas: An Integrated Proof Environment for Scalable Human–AI Collaborative Formalization*, 2026, [arXiv:2604.16347](#).
- [4] A. D. Brucker and B. Wolff, *Isabelle/DOF: Design and Implementation*, in *Software Engineering and Formal Methods*, 2019, pp. 275–292.
- [5] T. Kiecker et al., *CASCADE: Detecting Inconsistencies between Code and Documentation with Automatic Test Generation*, 2026, [arXiv:2604.19400](#).
- [6] R. A. DeMillo, R. J. Lipton, and F. G. Sayward, *Hints on test data selection: Help for the practicing programmer*, *IEEE Computer* 11(4), 1978, pp. 34–41.
- [7] Y. Jia and M. Harman, *An analysis and survey of the development of mutation testing*, *IEEE Transactions on Software Engineering* 37(5), 2011, pp. 649–678.
- [8] O. C. Z. Gotel and A. C. W. Finkelstein, *An analysis of the requirements traceability problem*, in *Proc. IEEE International Conference on Requirements Engineering*, 1994, pp. 94–101.
- [9] T. Kelly and R. Weaver, *The Goal Structuring Notation: a safety argument notation*, in *Proc. DSN Workshop on Assurance Cases*, 2004.